

Homework 3: Lucas - Kanade Tracking

Srividya Gandikota
 sgandiko
 sgandiko@andrew.cmu.edu

1 Lucas-Kanade Tracking

- 1.1 Starting with an initial guess of p (for instance, $p = [0, 0]^T$), we can compute the optimal p^* iteratively. In each iteration, the objective function is locally linearized by first-order Taylor expansion,

$$I_{t+1}(x' + \Delta p) \approx I_{t+1}(x') + \frac{\partial I_{t+1}(x')}{\partial x'^T} \frac{\partial W(x; p)}{\partial p^T} \Delta p \quad (1)$$

where $\Delta p = [\Delta p_x, \Delta p_y]^T$ is the template offset. Further, $x' = W(x; p) = x + p$ and $\frac{\partial I(x')}{\partial x'^T}$ is a vector of the x and y image gradients at pixel coordinate x' . In a similar manner to Equation 3 one can incorporate these linearized approximations into a vectorized form such that,

$$\arg \max_{\Delta p} \|A\Delta p - b\|_2^2 \quad (2)$$

such that $p \leftarrow p + \Delta p$ at each iteration.

1. What is $\frac{\partial W(x; p)}{\partial p^T}$? Write out the mathematical expression

- Derivation Below:

$$W(x; p) = \begin{bmatrix} 1 & 0 & p_1 \\ 0 & 1 & p_2 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + p_1 \\ y + p_2 \end{bmatrix}$$

The derivative of this matrix is:

$$\frac{\partial W(x; p)}{\partial p^T} = \begin{bmatrix} \frac{\partial W_x(x; p)}{\partial p_1^T} & \frac{\partial W_x(x; p)}{\partial p_2^T} \\ \frac{\partial W_y(x; p)}{\partial p_1^T} & \frac{\partial W_y(x; p)}{\partial p_2^T} \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

2. What is A and b ? Write out the mathematical expression

- A ($m \times n$) will become matrix of: $\frac{\partial I_{t+1}(x')}{\partial x'^T} \frac{\partial W(x; p)}{\partial p^T}$

$$A = \begin{bmatrix} \frac{\partial I_{t+1}(x'_1)}{\partial x'_1 T} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \frac{\partial I_{t+1}(x'_n)}{\partial x'_n T} \end{bmatrix} \begin{bmatrix} \frac{\partial W(x_1; p)}{\partial p^T} \\ \vdots \\ \frac{\partial W(x_n; p)}{\partial p^T} \end{bmatrix}$$

- b will become matrix of: $I_t(x) - I_{t+1}(x')$

$$b = \begin{bmatrix} I_t(x_1) - I_{t+1}(x'_1) \\ \vdots \\ I_t(x_n) - I_{t+1}(x'_n) \end{bmatrix}$$

3. What conditions must $A^T A$ meet so that a unique solution to Δp can be found?

- $A^T A$ must be both invertible and have a non zero determinant for a unique solution to be found. These requirements come from the need to find the inverse of the Hessian, which demands that the Hessian is invertible.

1.2 Implement a function with the following signature `LucasKanade(It, It1, rect, p0 = np.zeros(2))` that computes the optimal local motion from frame `It` to frame `It+1` that minimizes Equation 3.

My implementation for this function can be found in Figure 1 below.

```

1 import numpy as np
2 from scipy.interpolate import RectBivariateSpline
3
4 def LucasKanade(It, It1, rect, threshold, num_iters, p0=np.zeros(2)):
5     """
6         :param It: template image
7         :param It1: Current image
8         :param rect: Current position of the car (top left, bot right coordinates)
9         :param threshold: if the length of dp is smaller than the threshold, terminate the optimization
10        :param num_iters: number of iterations of the optimization
11        :param p0: Initial movement vector [dp_x0, dp_y0]
12        :return: p: movement vector [dp_x, dp_y]
13    """
14
15    # Put your implementation here
16    p = p0
17    magnitude = 1
18    count = 1
19
20    #Get Image Shapes
21    template_frame_height, template_frame_width = It.shape
22    init_frame_height, init_frame_width = It1.shape
23
24    #Set Up Rect Info
25    x_1, y_1, x_2, y_2 = rect
26    rect_height = int(x_2 - x_1)
27    rect_width = int(y_2 - y_1)
28
29    #Calculate Splines
30    template_x = np.linspace(0, template_frame_height, template_frame_height, False)
31    template_y = np.linspace(0, template_frame_width, template_frame_width, False)
32    template_spline = RectBivariateSpline(template_x, template_y, It)
33
34    init_x = np.linspace(0, init_frame_height, init_frame_height, False)
35    init_y = np.linspace(0, init_frame_width, init_frame_width, False)
36    init_spline = RectBivariateSpline(init_x, init_y, It1)
37
38    # Creating a 2D grid of coordinates (x, y)
39    complex_num = 1j
40    x, y = np.mgrid[x_1:(x_2+1):(rect_width * complex_num), y_1:(y_2+1):(rect_height * complex_num)]
41
42    while (magnitude > threshold) and (count < num_iters):
43        #Partial Derivatives
44        movement_x = x + p[0]
45        movement_y = y + p[1]
46        dp_x = init_spline.ev(movement_y, movement_x, dy = 1).flatten()
47        dp_y = init_spline.ev(movement_y, movement_x, dx = 1).flatten()
48
49        It1_p = init_spline.ev(movement_y, movement_x).flatten()
50        It_p = template_spline.ev(y, x).flatten()
51
52        #Make Matrix A
53        A = np.zeros((rect_width * rect_height, 2 * rect_width * rect_height))
54        for i in range(rect_width * rect_height):
55            A[i, (2 * i)] = dp_x[i]
56            A[i, (2 * i + 1)] = dp_y[i]
57        A = np.dot(A, (np.matlib.repmat(np.eye(2), rect_width * rect_height, 1)))
58
59        #Make Matrix b
60        b = np.reshape(It_p - It1_p, (rect_width * rect_height, 1))
61
62        #Pseudoinverse Solve and Update
63        magnitude = np.linalg.norm(np.linalg.pinv(A).dot(b))
64
65        #New p Update
66        p = (p + np.linalg.pinv(A).dot(b).T).ravel()
67        count += 1
68
69    return p

```

Figure 1: LucasKanade Function Implementation

1.3 Write a script testCarSequence.py that loads the video frames from carseq.npy, and runs the Lucas-Kanade tracker that you have implemented in the previous task to track the car. Similarly, write a script testGirlSequence.py that loads the video from girlseq.npy and runs your Lucas-Kanade tracker on it.

My implementation for the testcarsequence function can be found in Figure 2 below. The results of the

```

1 import argparse
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import matplotlib.patches as patches
5
6 # write your script here, we recommend the above libraries for making your animation
7 import LucasKanade
8
9 parser = argparse.ArgumentParser()
10 parser.add_argument('--num_iters', type=int, default=1e4, help='number of iterations of Lucas-Kanade')
11 parser.add_argument('--threshold', type=float, default=1e-2, help='dp threshold of Lucas-Kanade for terminating optimization')
12 args = parser.parse_args()
13 num_iters = args.num_iters
14 threshold = args.threshold
15
16 seq = np.load("../data/carseq.npy")
17 rect = [59, 116, 145, 151]
18
19 #Set up array for all boxes
20 _, _ , frame = seq.shape
21 box_array = rect.copy()
22 for i in range(frame - 1):
23     print(f'Frame {(i+1)}, of {frame}')
24
25     #Get frames for analysis
26     template_frame = seq[:, :, i]
27     curr_frame = seq[:, :, i + 1]
28
29     #Calculate Lucas Kanade Update
30     movement_vec = LucasKanade.LucasKanade(template_frame, curr_frame, rect, threshold, num_iters)
31     #print(movement_vec)
32     for j in range(4):
33         rect[j] += movement_vec[j % 2]
34     box_array = np.vstack((box_array, rect))
35
36     #Create Boxed Image
37     if i == 0 or i == 99 or i == 199 or i == 299 or i == 399:
38         print("True")
39         plt.figure()
40         plt.imshow(curr_frame, cmap='gray')
41         patch = patches.Rectangle((rect[0], rect[1]), (rect[2] - rect[0]),
42                                   (rect[3] - rect[1]), edgecolor='r', linewidth=3, facecolor='none')
43         ax = plt.gca()
44         ax.add_patch(patch)
45         plt.savefig('../results/Frame_'+str(i+1)+'.png', bbox_inches='tight')
46
47 np.save('carseqrects.npy', box_array)

```

Figure 2: testCarSequence Function Implementation

Lucas-Kanade tracker on the car can be found in Figure 3 below. Here we can see that the bounding box rectangle follows the car fairly well but as the car moves more, the box hangs off the edge of the car instead of accurately bounding against the entire car.

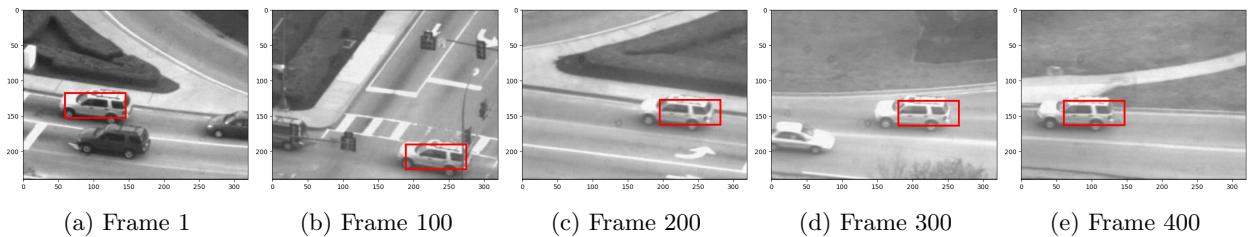


Figure 3: Lucas-Kanade Tracking on Car with One Single Template

My implementation for the testgirlsequence function can be found in Figure 4 below.

```

1 import argparse
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import matplotlib.patches as patches
5
6 # write your script here, we recommend the above libraries for making your animation
7 import LucasKanade
8
9 parser = argparse.ArgumentParser()
10 parser.add_argument('--num_iters', type=int, default=1e4, help='number of iterations of Lucas-Kanade')
11 parser.add_argument('--threshold', type=float, default=1e-2, help='dp threshold of Lucas-Kanade for terminating optimization')
12 args = parser.parse_args()
13 num_iters = args.num_iters
14 threshold = args.threshold
15
16 seq = np.load("../data/girlseq.npy")
17 rect = [280, 152, 330, 318]
18
19 #Set up array for all boxes
20 _, _, frame = seq.shape
21 box_array = rect.copy()
22 for i in range(frame - 1):
23     print(f'Frame {i+1}, of {frame}')
24
25 #Get frames for analysis
26 template_frame = seq[:, :, i]
27 curr_frame = seq[:, :, i + 1]
28
29 #Calculate Lucas Kanade Update
30 movement_vec = LucasKanade.LucasKanade(template_frame, curr_frame, rect, threshold, num_iters)
31 #print(movement_vec)
32 for j in range(4):
33     rect[j] += movement_vec[j % 2]
34 box_array = np.vstack((box_array, rect))
35
36 #Create Boxed Image
37 if i == 0 or i == 19 or i == 39 or i == 59 or i == 79:
38     print("True")
39     plt.figure()
40     plt.imshow(curr_frame, cmap='gray')
41     patch = patches.Rectangle((rect[0], rect[1]), (rect[2] - rect[0]),
42                               (rect[3] - rect[1]), edgecolor='r', linewidth=3, facecolor='none')
43     ax = plt.gca()
44     ax.add_patch(patch)
45     plt.savefig('../results/GirlFrame_'+str(i+1)+'.png', bbox_inches='tight')
46
47 np.save('girlseqrects.npy', box_array)

```

Figure 4: testGirlSequence Function Implementation

The results of the Lucas-Kanade tracker on the girl can be found in Figure 5 below. Here we can see that the bounding box rectangle follows the girl okay but as the girl moves more in front of more people, the box struggles to properly find the girl with the gradient difference. Once the girl leaves the region of other people, we notice that the box gets better at identifying her but still is not perfect.

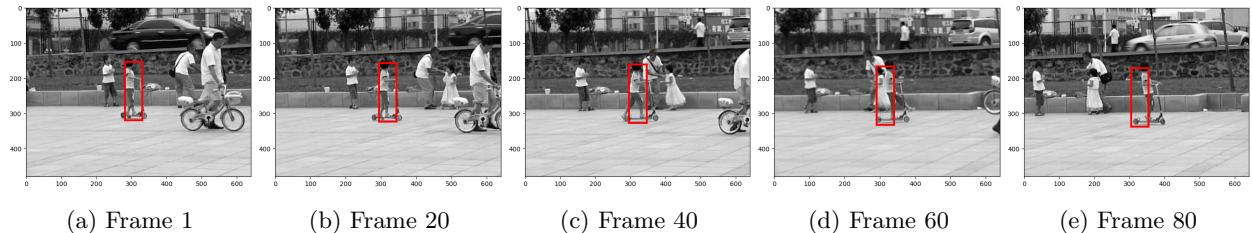


Figure 5: Lucas-Kanade Tracking on Girl with One Single Template

1.4 Write two scripts with a similar functionality to Q1.3 but with a template correction routine incorporated: testCarSequenceWithTemplateCorrection.py and testGirlSequenceWithTemplateCorrection.py. Save the rects as carseqrects-wcrt.npy and girlseqrects-wcrt.npy, and also report the performance at those frames.

My implementation for the testCarSequenceWithTemplateCorrection function can be found in Figure 6 below.

```

1 import argparse
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import matplotlib.patches as patches
5 import LucasKanade
6
7 # write your script here, we recommend the above libraries for making your animation
8
9 parser = argparse.ArgumentParser()
10 parser.add_argument('--num_iters', type=int, default=104, help='number of iterations of Lucas-Kanade')
11 parser.add_argument('--threshold', type=float, default=2, help='dp threshold of Lucas-Kanade for terminating optimization')
12 parser.add_argument('--template_threshold', type=float, default=5, help='threshold for determining whether to update template')
13 args = parser.parse_args()
14 num_iters = args.num_iters
15 threshold = args.threshold
16 template_threshold = args.template_threshold
17
18 seq = np.load('../data/carseq.npy')
19 rect = [59, 116, 145, 151]
20
21 #Set up array for all boxes
22 _, _, frame = seq.shape
23 box_array = rect.copy()
24 template_array = rect.copy()
25 template_array[0] = np.zeros(2)
26 move_vec_orig = np.zeros(2)
27
28 for i in range(frame-1):
29     print(f'Frame {i+1}, of {frame}')
30     #Get frames for analysis
31     curr_frame = seq[:, :, i+1]
32
33     #Calculate Lucas Kanade Update with Drift Correction
34     movement_vec = LucasKanade.LucasKanade(template_array, curr_frame, template_array, threshold, num_iters, move_vec_orig)
35     adjusted_vec = movement_vec + [template_array[0]-rect[0], template_array[1]-rect[1]]
36     updated_vec = LucasKanade.LucasKanade(seq[:, :, i], curr_frame, rect, threshold, num_iters, adjusted_vec)
37     delta_movement = np.linalg.norm(adjusted_vec - updated_vec)
38     #print(delta_movement)
39
40     if delta_movement < template_threshold:
41         change_vec = (updated_vec - [template_array[0]-rect[0], template_array[1]-rect[1]])
42         for j in range(4):
43             template_array[j] += change_vec[j % 2]
44             template_array = seq[:, :, i+1]
45             box_array = np.vstack((box_array, template_array))
46             move_vec_orig = np.zeros(2)
47     else:
48         box_array = np.vstack(
49             (box_array, [template_array[0]+movement_vec[0], template_array[1]+movement_vec[1], template_array[2]+movement_vec[0], template_array[3]+movement_vec[1]]))
50         move_vec_orig = movement_vec
51     #print(box_array)
52
53     #Create Boxed Image
54     carseq = np.load('carseqrects.npy')
55     if i == 0 or i == 99 or i == 199 or i == 299 or i == 399:
56         print(f'{i} frame")
57         plt.figure()
58         plt.imshow(curr_frame, cmap='gray')
59         no_update = carseq[i+1, :]
60         updated = box_array[i, :]
61         patch_no_update = patches.Rectangle((no_update[0], no_update[1]), (no_update[2] - no_update[0]),
62                                             (no_update[3] - no_update[1]), edgecolor='b', linewidth=3, facecolor='none')
63         patch_update = patches.Rectangle(updated[0], updated[1], (updated[2] - updated[0]),
64                                         (updated[3] - updated[1]), edgecolor='r', linewidth=3, facecolor='none')
65         ax = plt.gca()
66         ax.add_patch(patch_no_update)
67         ax.add_patch(patch_update)
68         plt.savefig('../results/CarCorrectionFrame_' + str(i+1) + '.png', bbox_inches='tight')
69
70 np.save('carseqrects-wcrt.npy', box_array)

```

Figure 6: testCarSequenceWithTemplateCorrection Function Implementation

The results of the Lucas-Kanade tracker on the car can be found in Figure 7 below. Here we can see that the blue no update template bounding box follows the car fairly well but as the car moves more, it lags behind and doesn't track the car well. The red bounding box uses the template correction method and we can see here that the box follows the car very well and follows the car much better than the no update version.

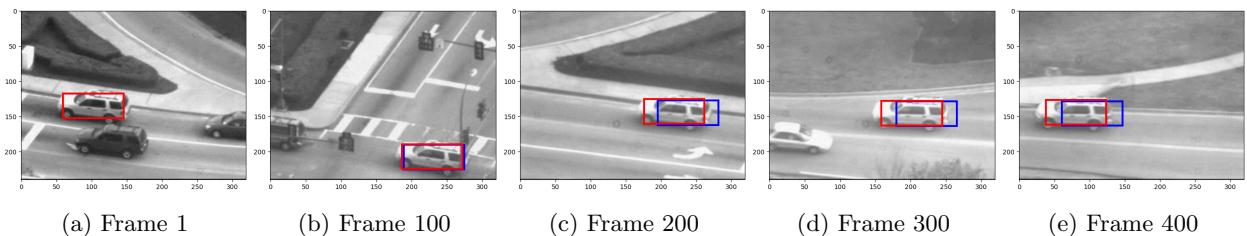


Figure 7: Lucas-Kanade Tracking on Car with Template Correction

My implementation for the testGirlSequenceWithTemplateCorrection function can be found in Figure 8 below.

```

1 import argparse
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import matplotlib.patches as patches
5 import LucasKanade
6
7 # write your script here, we recommend the above libraries for making your animation
8
9 parser = argparse.ArgumentParser()
10 parser.add_argument('--num_iters', type=int, default=4, help='number of iterations of Lucas-Kanade')
11 parser.add_argument('--threshold', type=float, default=1e-2, help='dp threshold of Lucas-Kanade for terminating optimization')
12 parser.add_argument('--template_threshold', type=float, default=5, help='threshold for determining whether to update template')
13 args = parser.parse_args()
14 num_iters = args.num_iters
15 threshold = args.threshold
16 template_threshold = args.template_threshold
17
18 seq = np.load("../data/girlseq.npy")
19 rect = [280, 152, 330, 318]
20
21 #Set up array for all boxes
22 _, _ = frame = seq.shape
23 print(seq.shape)
24 box_array = rect.copy()
25 template_array = rect.copy()
26 template_frame = seq[:, :, 0]
27 move_vec_orig = np.zeros(2)
28
29 for i in range(frame - 1):
30     print(f'Frame {i+1} of {frame}')
31     #Get frames for analysis
32     curr_frame = seq[:, :, i+1]
33
34     #Calculate Lucas Kanade Update with Drift Correction
35     movement_vec = LucasKanade.LucasKanade(template_frame, curr_frame, template_array, threshold, num_iters, move_vec_orig)
36     adjusted_vec = movement_vec + [template_array[0] - rect[0], template_array[1] - rect[1]]
37     updated_vec = LucasKanade.LucasKanade(seq[:, :, 0], curr_frame, rect, threshold, num_iters, adjusted_vec)
38     delta_movement = np.linalg.norm(adjusted_vec - updated_vec)
39     print(delta_movement)
40
41     if delta_movement < template_threshold:
42         change_vec = (updated_vec - [template_array[0]-rect[0], template_array[1]-rect[1]])
43         for j in range(4):
44             box_array[j] += change_vec[j % 2]
45             template_frame = seq[:, :, i+1]
46             box_array = np.vstack((box_array, template_array))
47             move_vec_orig = np.zeros(2)
48     else:
49         box_array = np.vstack([
50             [box_array, [template_array[0]+movement_vec[0], template_array[1]+movement_vec[1], template_array[2]+movement_vec[0], template_array[3]+movement_vec[1]]])
51         move_vec_orig = movement_vec
52     print(box_array)
53
54     #Create Boxed Image
55     girlseq = np.load('girlseqrects.npy')
56     if i == 0 or i == 19 or i == 39 or i == 59 or i == 79:
57         print("True")
58         plt.figure()
59         plt.imshow(curr_frame, cmap='gray')
60         no_update = girlseq[i+1, :]
61         updated = box_array[i, :]
62         patch_no_update = patches.Rectangle((no_update[0], no_update[1]), (no_update[2] - no_update[0]),
63                                           (no_update[3] - no_update[1]), edgecolor='b', linewidth=3, facecolor='none')
64         patch_update = patches.Rectangle(updated[0], updated[1], (updated[2] - updated[0]),
65                                         (updated[3] - updated[1]), edgecolor='r', linewidth=3, facecolor='none')
66         ax = plt.gca()
67         ax.add_patch(patch_no_update)
68         ax.add_patch(patch_update)
69         plt.savefig('../results/GirlCorrectionFrame' + str(i+1)+'.png', bbox_inches='tight')
70
71 np.save('girlseqrects-wcrt.npy', box_array)

```

Figure 8: testGirlSequenceWithTemplateCorrection Function Implementation

The results of the Lucas-Kanade tracker on the girl can be found in Figure 9 below. Here we can see that the blue no update template bounding box follows the girl okay but as the girl moves more in front of more people, the box struggles to properly find the girl with the gradient difference. Once the girl leaves the region of other people, we notice that the box gets better at identifying her but still is not perfect. On the other hand, the red template correction bounding box can be seen to follow the child perfectly even while it crosses the other background people.

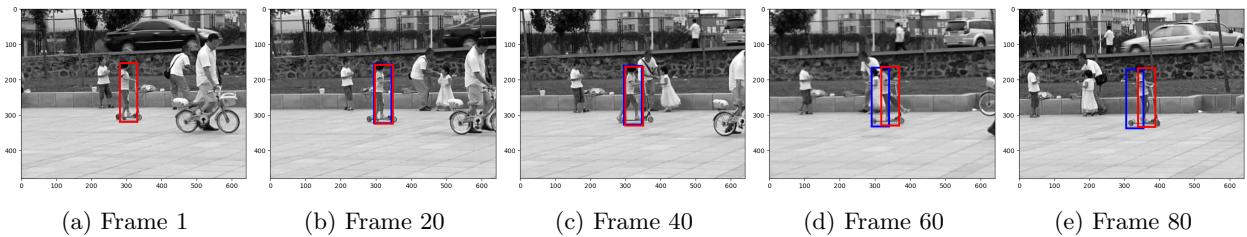


Figure 9: Lucas-Kanade Tracking on Girl with Template Correction

2 Affine Motion Subtraction

- 2.1 Write a function with the following signature `LucasKanadeAffine(It, It1)` which returns the affine transformation matrix M (shape 3×3), and It and $It1$ are It and $It+1$ respectively. `LucasKanadeAffine` should be relatively similar to `LucasKanade` from the first section (you will probably also find `scipy.ndimage.affine` transform helpful).

My implementation for the `LucasKanadeAffine` function can be found in Figure 10 below.

```

1 import numpy as np
2 from scipy.interpolate import RectBivariateSpline
3
4 def LucasKanadeAffine(It, It1, threshold, num_iters):
5     """
6         :param It: template image
7         :param It1: Current image
8         :param threshold: if the length of dp is smaller than the threshold, terminate the optimization
9         :param num_iters: number of iterations of the optimization
10        :return: M: the Affine warp matrix [2x3 numpy array] put your implementation here
11    """
12    M = np.array([[1.0, 0.0, 0.0], [0.0, 1.0, 0.0]])
13    p = np.zeros(6)
14    error = 1
15    count = 1
16
17    #Get Image Shapes
18    template_frame_height, template_frame_width = It.shape
19    init_frame_height, init_frame_width = It1.shape
20
21    #Calculate Splines
22    template_x = np.linspace(0, template_frame_height, template_frame_height, False)
23    template_y = np.linspace(0, template_frame_width, template_frame_width, False)
24    template_spline = RectBivariateSpline(template_x, template_y, It)
25
26    init_x = np.linspace(0, init_frame_height, init_frame_height, False)
27    init_y = np.linspace(0, init_frame_width, init_frame_width, False)
28    init_spline = RectBivariateSpline(init_x, init_y, It1)
29
30    #Creating a mesh grid of coordinates (x, y)
31    grid_x, grid_y = np.meshgrid(range(template_frame_width), range(template_frame_height))
32    x = grid_x.flatten()
33    y = grid_y.flatten()
34    coords = np.vstack((x, y, np.ones(template_frame_height * template_frame_width)))
35
36    while error > threshold and count < num_iters:
37        count += 1
38        #Update transformation matrix M with p
39        M[0, 0] = 1 + p[0]
40        M[0, 1] = p[1]
41        M[0, 2] = p[2]
42        M[1, 0] = p[3]
43        M[1, 1] = p[4] + 1
44        M[1, 2] = p[5]
45
46        #Transform and remove out of bounds coordinates
47        warped_x = np.dot(M, coords)[0]
48        warped_y = np.dot(M, coords)[1]
49        bad_coords = find_bad_coordinates(warped_x, warped_y, template_frame_width, template_frame_height)
50        valid_coords = np.logical_not(bad_coords)
51        x_new = x[valid_coords]
52        x_new = x_new[:, np.newaxis]
53        y_new = y[valid_coords]
54        y_new = y_new[:, np.newaxis]
55        warped_x = warped_x[valid_coords]
56        warped_x = warped_x[:, np.newaxis]
57        warped_y = warped_y[valid_coords]
58        warped_y = warped_y[:, np.newaxis]
59
60        #Find partial derivatives
61        dp_x = init_spline.ev(warped_y, warped_x, dy = 1).flatten()
62        dp_x = dp_x[:, np.newaxis]
63        dp_y = init_spline.ev(warped_y, warped_x, dx = 1).flatten()
64        dp_y = dp_y[:, np.newaxis]
65        It_p = template_spline.ev(y_new.flatten(), x_new.flatten()).flatten()
66        It1_p = init_spline.ev(warped_y.flatten(), warped_x.flatten()).flatten()
67
68        #Calculate matrices
69        A = np.hstack((x_new * dp_x, y_new * dp_x, dp_x, x_new * dp_y, y_new * dp_y, dp_y))
70        b = (It_p - It1_p).reshape(-1, 1)
71
72        #PseudoInverse Solve and Update
73        magnitude = np.linalg.norm(np.linalg.pinv(A).dot(b))
74        p = (p + np.linalg.pinv(A).dot(b.T).ravel())

```

Figure 10: `LucasKanadeAffine` Function Implementation

- 2.2 Using the function you have developed for dominant motion estimation, write a function with the following signature SubtractDominantMotion(image1, image2) where image1 and image2 form the input image pair, and the return value mask is a binary image of the same size that dictates which pixels are considered to be corresponding to moving objects.

My implementation for the SubtractDominantMotion function can be found in Figure 11 below.

```
1 import numpy as np
2 import LucasKanadeAffine
3 import scipy
4
5 def SubtractDominantMotion(image1, image2, threshold, num_iters, tolerance):
6     """
7         :param image1: Images at time t
8         :param image2: Images at time t+1
9         :param threshold: used for LucasKanadeAffine
10        :param num_iters: used for LucasKanadeAffine
11        :param tolerance: binary threshold of intensity difference when computing the mask
12        :return: mask: [nxm]
13    """
14
15    # put your implementation here
16    mask = np.ones(image1.shape, dtype=bool)
17
18    #Find Matrix M and Warp Image
19    matrix_M = LucasKanadeAffine.LucasKanadeAffine(image1, image2, threshold, num_iters, tolerance)
20    warped_image = scipy.ndimage.affine_transform(image1, -matrix_M, offset=0.0, output_shape=None)
21
22    #Generate Image Mask
23    mask = (np.abs(warped_image - image2) > tolerance)
24    mask = scipy.ndimage.morphology.binary_erosion(mask)
25    mask = scipy.ndimage.morphology.binary_dilation(mask)
26
27    return mask
28
```

Figure 11: SubtractDominantMotion Function Implementation

2.3 Write two scripts `testAntSequence.py` and `testAerialSequence.py` that load the image sequence from `antseq.npy` and `aerialseq.npy` and run the motion detection routine you have developed to detect the moving objects. Try to implement `testAntSequence.py` first as it involves little camera movement and can help you debug your mask generation procedure.

My implementation for the `testAntSequence` function can be found in Figure 12 below.

```

1 import argparse
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import matplotlib.patches as patches
5 import SubtractDominantMotion
6
7 # write your script here, we recommend the above libraries for making your animation
8
9 parser = argparse.ArgumentParser()
10 parser.add_argument('--num_iters', type=int, default=1e3, help='number of iterations of Lucas-Kanade')
11 parser.add_argument('--threshold', type=float, default=1e-2, help='dp threshold of Lucas-Kanade for terminating optimization')
12 parser.add_argument('--tolerance', type=float, default=0.2, help='binary threshold of intensity difference when computing the mask')
13 args = parser.parse_args()
14 num_iters = args.num_iters
15 threshold = args.threshold
16 tolerance = args.tolerance
17
18 seq = np.load('../data/antseq.npy')
19 _, frame = seq.shape
20 for i in range(frame - 1):
21     print(f'Frame {i+1}, of {frame}')
22
23     #Get Frames for Analysis
24     template_frame = seq[:, :, i]
25     curr_frame = seq[:, :, i + 1]
26
27     #Get Image for Mask
28     mask = SubtractDominantMotion.SubtractDominantMotion(template_frame, curr_frame, threshold, num_iters, tolerance)
29     positions_x = []
30     positions_y = []
31     for j in range(mask.shape[0]):
32         for k in range(mask.shape[1]):
33             if mask[j, k] == 0:
34                 positions_y.append(j)
35                 positions_x.append(k)
36
37     #Create Boxed Image
38     if i == 29 or i == 59 or i == 89 or i == 119:
39         print("True")
40         plt.figure()
41         plt.imshow(curr_frame, cmap='gray')
42         plt.plot(positions_x, positions_y, '.', markerfacecolor='blue')
43         plt.savefig('../results/AntFrame_{str(i+1)}.png', bbox_inches='tight')
```

Figure 12: `testAntSequence` Function Implementation

The results of the Lucas-Kanade Affine tracker on the ant sequence can be found in Figure 13 below. Here we can see that the blue plots follows the the ants around fairly well even though it may be a little difficult to see. The affine transformation visibly helps the Lucas Kanade method follow the optical flow of the ants much better.

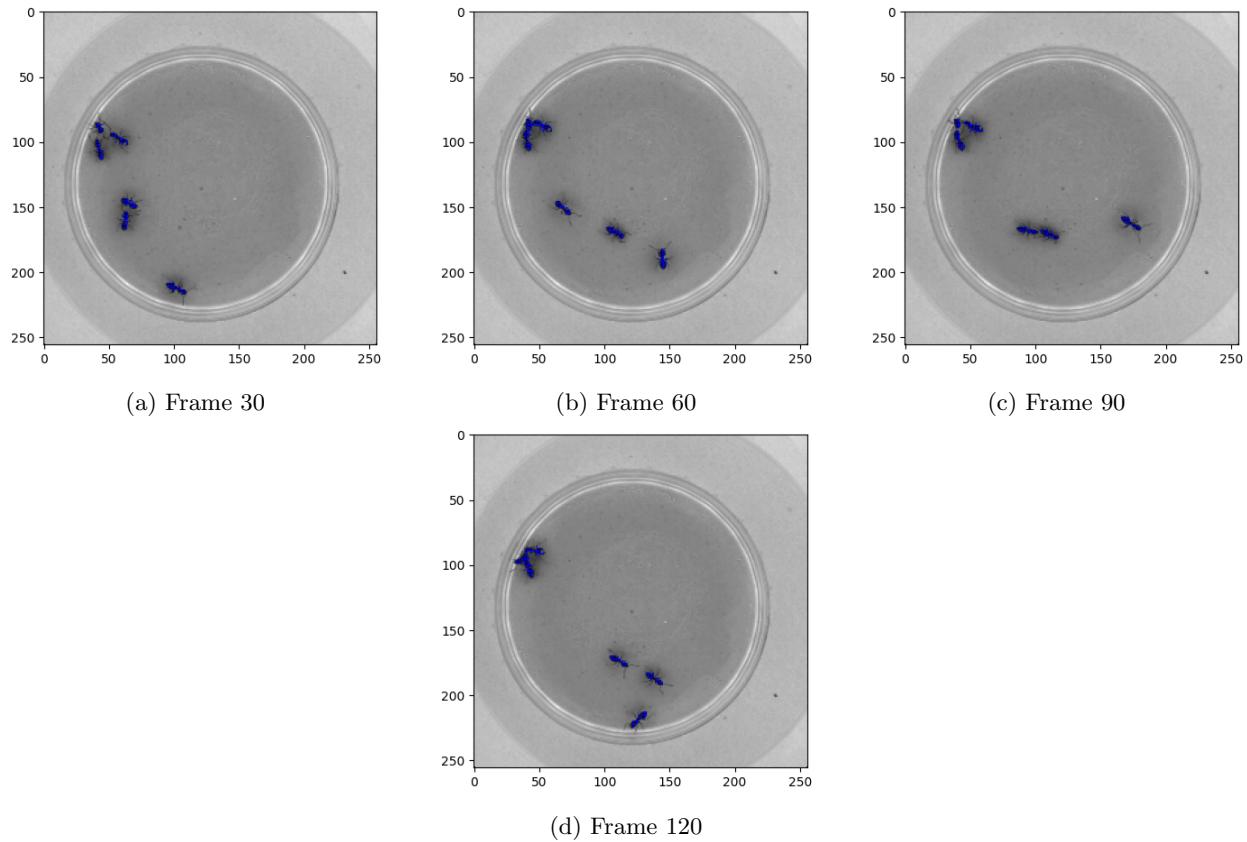


Figure 13: Lucas-Kanade Affine Tracking on Ants

My implementation for the `testAerialSequence` function can be found in Figure 14 below.

```

1 import argparse
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import matplotlib.patches as patches
5 import SubtractDominantMotion
6
7 # write your script here, we recommend the above libraries for making your animation
8
9 parser = argparse.ArgumentParser()
10 parser.add_argument('--num_iters', type=int, default=le3, help='number of iterations of Lucas-Kanade')
11 parser.add_argument('--threshold', type=float, default=le-2, help='dp threshold of Lucas-Kanade for terminating optimization')
12 parser.add_argument('--tolerance', type=float, default=0.2, help='binary threshold of intensity difference when computing the mask')
13 args = parser.parse_args()
14 num_iters = args.num_iters
15 threshold = args.threshold
16 tolerance = args.tolerance
17
18 seq = np.load('../data/aerialseq.npy')
19 _, _, frame = seq.shape
20 for i in range(frame - 1):
21     print(f'Frame {(i+1)} of {frame}')
22
23 #Get Frames for Analysis
24 template_frame = seq[:, :, i]
25 curr_frame = seq[:, :, i + 1]
26
27 #Get Image for Mask
28 mask = SubtractDominantMotion.SubtractDominantMotion(template_frame, curr_frame, threshold, num_iters, tolerance)
29 positions_x = []
30 positions_y = []
31 for j in range(mask.shape[0]):
32     for k in range(mask.shape[1]):
33         if mask[j, k] == 0:
34             positions_y.append(j)
35             positions_x.append(k)
36
37 #Create Boxed Image
38 if i == 29 or i == 59 or i == 89 or i == 119:
39     print("True")
40     plt.figure()
41     plt.imshow(curr_frame, cmap='gray')
42     plt.plot(positions_x, positions_y, ',', markerfacecolor='blue')
43     plt.savefig('../results/AerialFrame '+str(i+1)+'.png', bbox_inches='tight')

```

Figure 14: testAerialSequence Function Implementation

The results of the Lucas-Kanade Affine tracker on the aerial car sequence can be found in Figure 15 below. The blue outline in this image is much more difficult to see in these images but if we zoom in, we see that the blue tracks the car shadow to identify the locations by using the optical flow of the darker shadow regions instead.

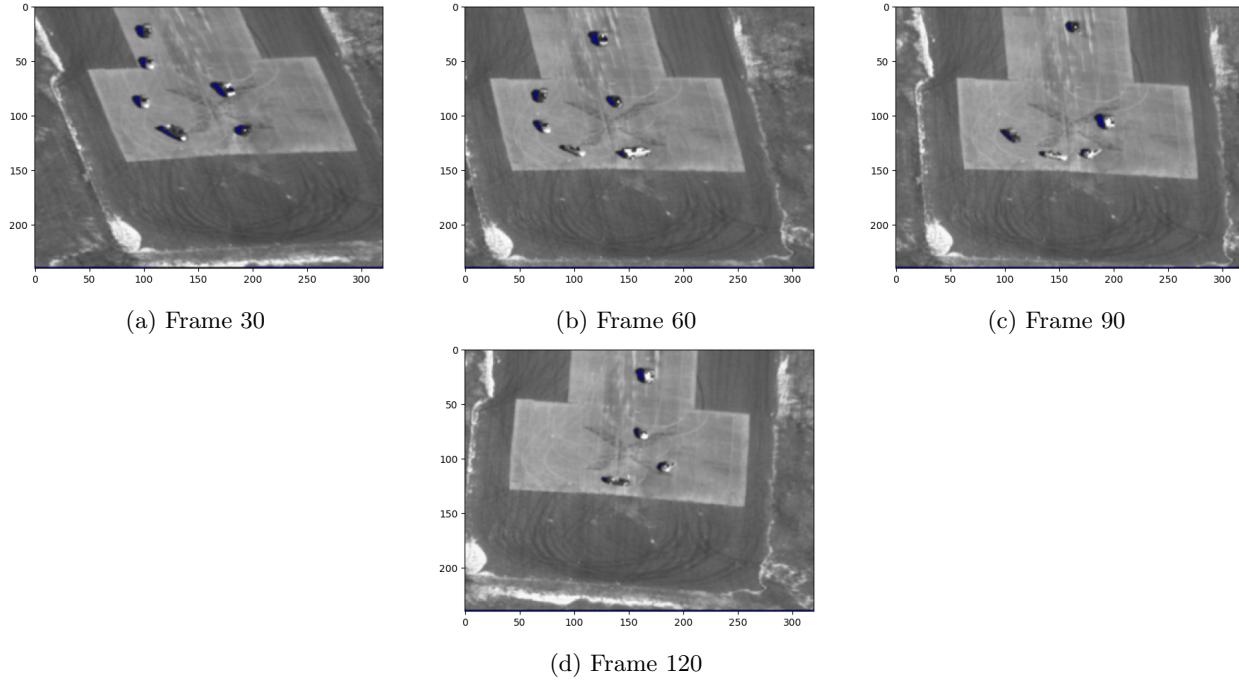


Figure 15: Lucas-Kanade Affine Tracking on Aerial Car Motion

3 Efficient Tracking

- 3.1 Reimplement the function LucasKanadeAffine(It,It1) as InverseCompositionAffine(It,It1) using the inverse compositional method. Report the results of the query frames (see Q2.3) for both ant and aerial sequences, as well as the runtime performance gain you observe using the inverse composition method. In your own words, please describe why the inverse compositional approach is more computationally efficient than the classical approach.

My implementation for the InverseCompositionAffine function can be found in Figure 16 below.

```
1 import numpy as np
2 from scipy.interpolate import RectBivariateSpline
3 import scipy
4
5 def InverseCompositionAffine(It, It1, threshold, num_iters):
6     """
7         :param It: template image
8         :param It1: Current image
9         :param threshold: if the length of dp is smaller than the threshold, terminate the optimization
10        :param num_iters: number of iterations of the optimization
11        :return: M: the Affine warp matrix [2x3 numpy array]
12    """
13
14    # put your implementation here
15    M = np.array([[1.0, 0.0, 0.0], [0.0, 1.0, 0.0]])
16    Hessian = np.zeros((6, 6))
17    p = np.zeros(6)
18    error = 1
19    count = 1
20
21    #Get Gradients and Calculate Hessian
22    x_grad, y_grad = np.gradient(It)
23    for i in range(It.shape[0]):
24        for j in range(It.shape[1]):
25            A_matrix = np.dot(np.array([x_grad[i, j], y_grad[i, j]]), np.array([[j, 0, i, 0, 1, 0], [0, j, 0, i, 0, 1]])[np.newaxis, :])
26            Hessian = Hessian + np.dot(A_matrix.T, A_matrix)
27    print("Hessian Found")
28
29    while error > threshold and count < num_iters:
30        print("Iteration: ", count)
31        count += 1
32        b_matrix = np.zeros((6, 1))
33        warped_curr_frame = scipy.ndimage.affine_transform(It1, M)
34
35        #Compute gradient of b
36        for i in range(It.shape[0]):
37            for j in range(It.shape[1]):
38                A_matrix = np.dot(np.array([x_grad[i, j], y_grad[i, j]]), np.array([[j, 0, i, 0, 1, 0], [0, j, 0, i, 0, 1]])[np.newaxis, :])
39                Hessian = Hessian + np.dot(A_matrix.T, A_matrix)
40                b_matrix = b_matrix + (np.transpose(A_matrix) * (warped_curr_frame - It))
41
42        #Pseudoinverse Solve and Update
43        magnitude = np.dot(np.linalg.pinv(Hessian), (b_matrix))
44
45        #Update M matrix using derivative
46        mag_derivative = [
47            [1.0 + magnitude[0][0], magnitude[2][0], magnitude[4][0]],
48            [magnitude[1][0], 1.0 + magnitude[3][0], magnitude[5][0]],
49            [0, 0, 1]]
50        M = np.dot(np.concatenate((M, np.array([[0, 0, 1]])), axis=0), np.linalg.pinv(mag_derivative))
51        M = M[0:2, :]
52
53        #New p Update
54        p = (p + magnitude.T).ravel()
55        error = np.linalg.norm(magnitude)
56
57    return M
```

Figure 16: InverseCompositionAffine Function Implementation

The results of the Lucas-Kanade Inverse tracker can be found in Figures 17 and 18 below. Here we can see that the blue follows the car and ant just as well as the affine method we saw before. The difference between the two is much more noticeable when considering the runtime between the two. For the original Lucas Kanade Affine method, the runtime for the ant and car were found to be 324 seconds (5.4 minutes) and 371 seconds (6.18 minutes) correspondingly. On the other hand, the inverse Lucas Kanade method was able to achieve the same results with a much faster runtime, 119 seconds (1.983 minutes) for ant and 158 seconds (2.63 minutes) for aerial car. This disparity is most likely due to the fact that the inverse method allows us to precompute the large terms like the hessian and gradients. Using this, we can track the template to the current frame and then use that to find the correct warp from the current frame to the template with inverted warping.

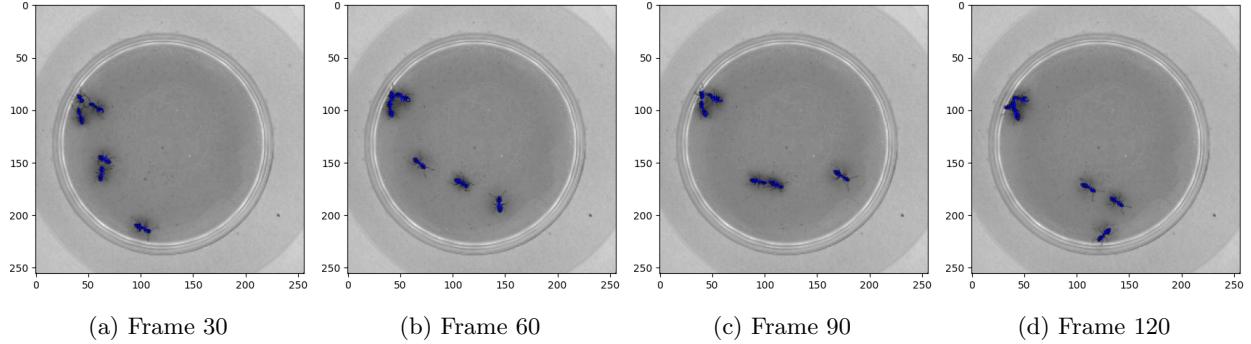


Figure 17: Lucas-Kanade Inverse Tracking on Ant Motion

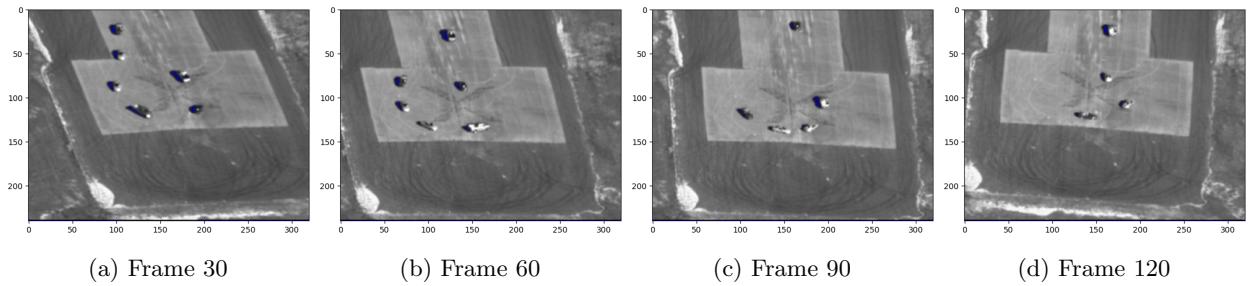


Figure 18: Lucas-Kanade Inverse Tracking on Aerial Car Motion